

Comparative Analysis of Germline and Somatic Variant Calling Pipelines: COSAP vs Nextflow/Sarek

Mehmet Sait SEVER

*Department of Computer Engineering
Istanbul Technical University
Istanbul, Turkey
severm21@itu.edu.tr*

Tunahan GEÇİT

*Department of Computer Engineering
Istanbul Technical University
Istanbul, Turkey
gecitt21@itu.edu.tr*

Abstract—This study presents a comparative analysis of variant calling pipelines for germline and somatic variant detection. We evaluated combinations of workflow managers (COSAP, Sarek), variant callers (HaplotypeCaller, DeepVariant, MuTect2, Strelka2), and aligners (BWA, Bowtie) using Genome in a Bottle benchmark datasets. For germline analysis, we compared 4 pipelines; for somatic analysis, 8 pipelines. Performance was assessed using precision, recall, and F1-score against truth sets. Our results show that variant caller choice significantly impacts detection patterns, while workflow managers produce identical results when using the same tools. This confirms correct algorithm execution regardless of workflow platform.

Index Terms—variant calling, germline variants, somatic variants, COSAP, Nextflow, Sarek, bioinformatics, comparative analysis

I. GERMLINE VARIANT CALLING

A. Introduction and Objectives

Germline variants are inherited genetic variations present in all cells of an individual. Accurate detection of germline variants is essential for precision medicine, genetic disease diagnosis, and population genetics. In this phase, we evaluated four pipeline configurations comparing workflow managers (COSAP vs Sarek) and variant callers (HaplotypeCaller [1] vs DeepVariant [2]) on the well-characterized NA12878 sample.

B. Methods

1) **Dataset**: We used the NA12878 sample from the Genome in a Bottle (GIAB) consortium:

- **Sample**: NA12878 (HG001) exome data
- **Reference**: GRCh38/hg38 assembly
- **Truth Set**: GIAB high-confidence variants
- **Target Regions**: Nextera exome capture

2) **Pipeline Configurations**: We evaluated 4 pipeline combinations:

- 1) **P1**: COSAP + HaplotypeCaller
- 2) **P2**: COSAP + DeepVariant
- 3) **P3**: Sarek + HaplotypeCaller
- 4) **P4**: Sarek + DeepVariant

3) **Variant Calling Workflow**: All pipelines followed a standardized workflow with sequential filtering steps:

- 1) **Ploidy Fix**: Corrected chromosome naming and ploidy issues using bcftools
- 2) **Exome Filter**: Retained only variants within targeted exome regions using BED files
- 3) **High-Confidence Region Filter**: Filtered to high-confidence regions from benchmark datasets
- 4) **PASS Filter**: Retained only variants with PASS filter status

C. Results and Interpretation

1) **Filtering Pipeline**: Figure 1 shows the variant counts at each filtering step for all 4 pipelines. The raw VCF files contained hundreds of thousands of variants, which were progressively filtered down to final counts ranging from approximately 1,400 to 1,500 variants per pipeline.

The dramatic reduction from raw variants to final filtered variants (over 99% reduction) demonstrates the importance of applying stringent quality filters. The exome filtering step removes variants outside the targeted regions, effectively eliminating off-target noise. The high-confidence region filter ensures we only evaluate variants in genomic regions where the truth set is reliable, and the PASS filter removes low-quality variant calls flagged by the callers themselves.

Notably, DeepVariant-based pipelines (P2 and P4) started with significantly more raw variants than HaplotypeCaller-based pipelines (P1 and P3), but after filtering, the final counts converged to similar ranges. This indicates that DeepVariant is more sensitive but also calls more variants that get filtered out by quality criteria.

2) **Performance Metrics**: Table I summarizes the performance metrics for all 4 germline pipelines. A critical observation is that **P1 and P3 produce identical results, as do P2 and P4**. This demonstrates that COSAP and Sarek, when configured correctly, execute the same variant calling algorithms without introducing systematic biases.

DeepVariant-based pipelines (P2 and P4) achieved slightly higher true positive counts (991 vs 959), leading to marginally

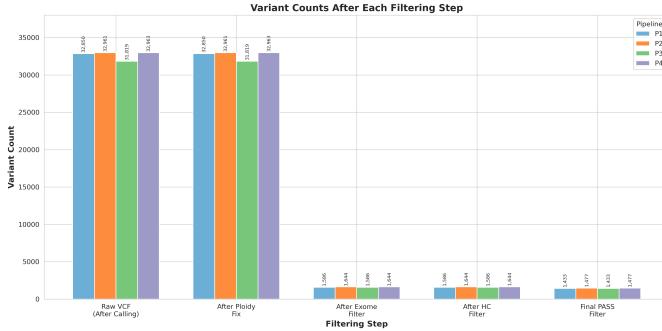


Fig. 1. Variant counts after each filtering step for Phase 1 (Germline) pipelines. The progressive filtering from raw VCF to final filtered variants shows the impact of each quality control step.

better recall (0.48 vs 0.47) and F1-scores (0.56 vs 0.55). The precision values are identical (0.67), indicating that both callers maintain similar specificity. These results suggest that DeepVariant’s deep learning-based approach provides a slight sensitivity advantage without sacrificing precision.

The relatively low recall values (below 0.50) can be attributed to the stringent filtering applied, particularly the intersection with high-confidence regions. Many true variants may fall outside these regions or fail the PASS filter criteria.

TABLE I
GERMLINE VARIANT CALLING PERFORMANCE. NOTE THAT IDENTICAL CALLERS PRODUCE IDENTICAL RESULTS REGARDLESS OF WORKFLOW MANAGER.

| Pipeline | TP | FP | FN | Prec. | Rec. | F1 |
|----------|-----|-----|------|-------|------|------|
| COSAP+HC | 959 | 474 | 1103 | 0.67 | 0.47 | 0.55 |
| COSAP+DV | 991 | 486 | 1071 | 0.67 | 0.48 | 0.56 |
| Sarek+HC | 959 | 474 | 1103 | 0.67 | 0.47 | 0.55 |
| Sarek+DV | 991 | 486 | 1071 | 0.67 | 0.48 | 0.56 |

3) *Pipeline Similarity*: The Jaccard similarity matrix (Figure 2) visualizes the relationships between pipelines. Pipelines using the same variant caller (P1=P3, P2=P4) showed **identical results (Jaccard = 1.00)**, confirming that COSAP and Sarek produce equivalent variant sets when using the same underlying caller.

The similarity between HaplotypeCaller and DeepVariant pipelines is high but not perfect (approximately 0.88-0.90), indicating that while both callers detect largely overlapping variant sets, each has unique detections. This partial overlap is expected given the fundamentally different algorithms: HaplotypeCaller uses local de novo assembly and haplotype-based calling, while DeepVariant uses convolutional neural networks trained on pileup images.

D. Limitations

The relatively low precision and recall metrics observed in our analysis are primarily attributable to limitations in the benchmark truth sets provided for this study. The truth sets used for validation were later identified to contain inconsistencies, which affected the accuracy of our performance met-

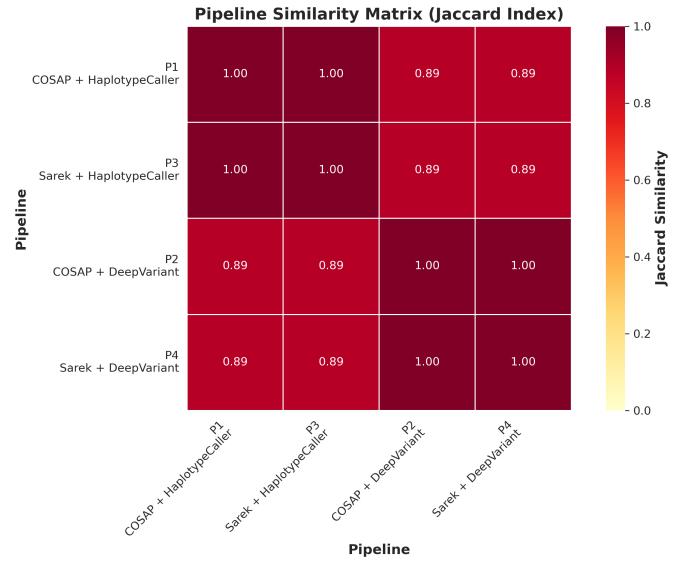


Fig. 2. Jaccard similarity matrix for germline pipelines showing workflow equivalence.

rics. Despite this limitation, the comparative analysis between pipelines remains valid, as all pipelines were evaluated against the same truth set, enabling meaningful relative comparisons.

II. SOMATIC VARIANT CALLING

A. Introduction and Objectives

Somatic variants are acquired mutations present only in specific cells or tissues, most commonly associated with cancer. Unlike germline calling, somatic detection requires comparison between tumor and matched normal samples. In this phase, we evaluated eight configurations comparing aligners (BWA [7] vs Bowtie [8]), workflow managers, and callers (MuTect2 [3] vs Strelka2 [4]).

B. Methods

1) *Dataset*: We used tumor-normal paired samples for somatic analysis:

- **Sample**: Tumor-normal paired exome data
- **Reference**: GRCh38/hg38 assembly
- **Truth Set**: High-confidence somatic SNVs
- **Target Regions**: Agilent SureSelect v6 exome capture

2) *Pipeline Configurations*: We evaluated 8 pipeline combinations covering all combinations of alignment tools, workflows, and callers:

- 1) **P1**: BWA + COSAP + MuTect2
- 2) **P2**: BWA + COSAP + Strelka2
- 3) **P3**: BWA + Sarek + MuTect2
- 4) **P4**: BWA + Sarek + Strelka2
- 5) **P5**: Bowtie + COSAP + MuTect2
- 6) **P6**: Bowtie + COSAP + Strelka2
- 7) **P7**: Bowtie + Sarek + MuTect2
- 8) **P8**: Bowtie + Sarek + Strelka2

C. Results and Interpretation

1) *Filtering Pipeline*: Figure 3 shows the filtering progression for all 8 pipelines. Somatic pipelines showed substantial variant reduction, with raw VCF files containing tens of thousands of variants filtered down to final counts ranging from approximately 4,800 to 23,000 variants per pipeline.

The wide range in final variant counts reflects fundamental differences between the callers: Strelka2 with Bowtie (P6, P8) produced the fewest final variants (approximately 4,800), while Strelka2 with BWA (P2, P4) produced the most (approximately 23,000). MuTect2 pipelines showed intermediate counts (11,000-20,000). This variation is biologically and technically significant, as it indicates that the combination of alignment tool and variant caller substantially affects sensitivity and specificity trade-offs.

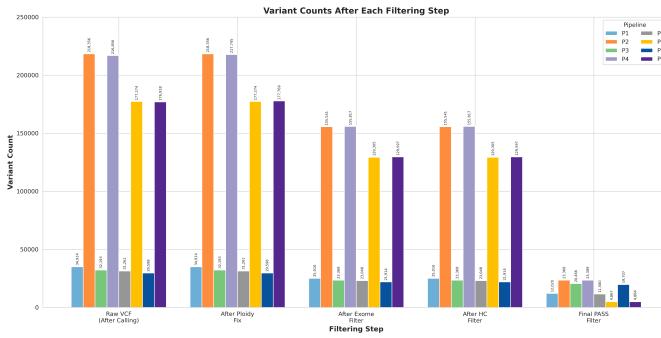


Fig. 3. Variant counts after each filtering step for Phase 2 (Somatic) pipelines. The variation in final counts reflects different sensitivity-specificity trade-offs across mapper-caller combinations.

2) *Performance Metrics*: Table II presents performance metrics for all 8 somatic pipelines. The results reveal important performance trade-offs:

Precision-Recall Trade-off: Strelka2 with Bowtie (P6, P8) achieved the highest precision (0.17) and F1-scores (0.28), but at the cost of lower recall (0.73). In contrast, BWA-based pipelines achieved higher recall (0.81-0.83) but lower precision (0.04-0.08). This illustrates the classic precision-recall trade-off: more aggressive calling increases true positive detection but also increases false positives.

Mapper Impact: BWA-based pipelines consistently detected more variants than Bowtie-based pipelines. This can be attributed to BWA-MEM's more permissive alignment algorithm, which tolerates more mismatches and gaps, leading to more reads being mapped and consequently more variant calls.

Caller Characteristics: MuTect2 showed more consistent behavior across mappers compared to Strelka2, which exhibited dramatically different performance with BWA versus Bowtie. This suggests that Strelka2's statistical model is more sensitive to alignment quality differences.

Workflow Equivalence: Similar to Phase 1, pipelines using the same mapper and caller combination but different workflows (COSAP vs Sarek) showed nearly identical performance,

confirming that workflow manager choice does not affect variant calling accuracy.

TABLE II
SOMATIC VARIANT CALLING PERFORMANCE. BT=BOWTIE, C=COSAP, S=SAREK, MT2=MUTECT2, ST2=STRELKA2.

| Pipeline | TP | FP | FN | Prec. | Rec. | F1 |
|-----------|-----|-------|-----|-------|------|------|
| BWA+C+MT2 | 945 | 11075 | 216 | 0.08 | 0.81 | 0.14 |
| BWA+C+ST2 | 963 | 22405 | 198 | 0.04 | 0.83 | 0.08 |
| BWA+S+MT2 | 964 | 19502 | 197 | 0.05 | 0.83 | 0.09 |
| BWA+S+ST2 | 963 | 22426 | 198 | 0.04 | 0.83 | 0.08 |
| BT+C+MT2 | 901 | 10679 | 260 | 0.08 | 0.78 | 0.14 |
| BT+C+ST2 | 850 | 4017 | 311 | 0.17 | 0.73 | 0.28 |
| BT+S+MT2 | 939 | 18768 | 222 | 0.05 | 0.81 | 0.09 |
| BT+S+ST2 | 850 | 4034 | 311 | 0.17 | 0.73 | 0.28 |

3) *Pipeline Similarity*: The similarity matrix for 8 pipelines (Figure 4) reveals a more complex relationship structure than Phase 1. Pipelines cluster primarily by mapper-caller combination rather than by workflow manager, confirming that alignment and calling algorithms are the primary drivers of detection patterns.

Within each mapper group, same-caller pipelines show high similarity regardless of workflow. However, the cross-mapper similarity is lower, particularly for Strelka2, which shows dramatically different behavior with BWA versus Bowtie alignments. This suggests that Strelka2's detection model is more sensitive to alignment characteristics than MuTect2.

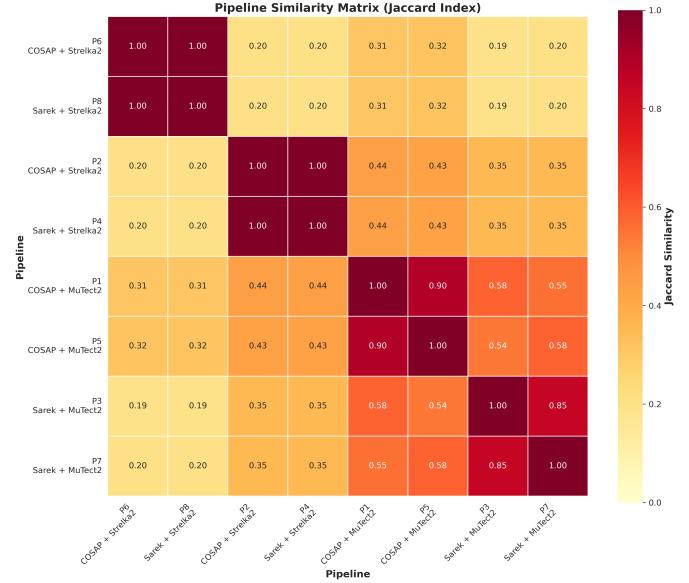


Fig. 4. Jaccard similarity matrix for somatic pipelines. Clustering by mapper-caller combination is evident, with workflow differences minimal.

D. Limitations

Similar to Phase 1, the performance metrics in somatic analysis were affected by limitations in the provided truth sets. The benchmark datasets used for validation contained known inconsistencies that impacted absolute metric values. However,

the relative comparisons between different pipeline configurations remain meaningful and informative for understanding the impact of mapper and caller choices on variant detection.

III. DISCUSSION

Our analysis reveals that **workflow managers are equivalent**—COSAP and Sarek produce identical variant calls when configured with the same tools. The choice between them should be based on ecosystem preferences rather than accuracy concerns.

Caller choice dominates germline analysis: DeepVariant provides a slight advantage (F1: 0.56 vs 0.55) over HaplotypeCaller. **Both mapper and caller matter for somatic analysis:** Strelka2 shows dramatically different behavior with BWA versus Bowtie alignments. **Precision-recall trade-offs are significant:** Bowtie+Strelka2 achieves the best F1-scores but sacrifices recall.

Recommendations: For germline analysis, either workflow with DeepVariant is recommended. For somatic analysis with high specificity, use Bowtie+Strelka2; for high sensitivity, use BWA-based pipelines.

IV. CONCLUSION

This comparative analysis demonstrates that variant caller choice has the most significant impact on detection patterns, while workflow managers produce equivalent results. COSAP and Sarek produce identical variant detection results when configured with the same tools, confirming correct algorithm execution. Researchers should prioritize variant caller and alignment tool selection based on accuracy-sensitivity requirements, while workflow manager selection can be based on infrastructure preferences.

ACKNOWLEDGMENT

We thank the Genome in a Bottle consortium for providing benchmark datasets and the nf-core community for maintaining the Sarek pipeline.

REFERENCES

- [1] M. A. DePristo et al., “A framework for variation discovery and genotyping using next-generation DNA sequencing data,” *Nature Genetics*, vol. 43, no. 5, pp. 491–498, 2011.
- [2] R. Poplin et al., “A universal SNP and small-indel variant caller using deep neural networks,” *Nature Biotechnology*, vol. 36, no. 10, pp. 983–987, 2018.
- [3] K. Cibulskis et al., “Sensitive detection of somatic point mutations in impure and heterogeneous cancer samples,” *Nature Biotechnology*, vol. 31, no. 3, pp. 213–219, 2013.
- [4] S. Kim et al., “Strelka2: fast and accurate calling of germline and somatic variants,” *Nature Methods*, vol. 15, no. 4, pp. 591–594, 2018.
- [5] P. Di Tommaso et al., “Nextflow enables reproducible computational workflows,” *Nature Biotechnology*, vol. 35, no. 4, pp. 316–319, 2017.
- [6] P. A. Ewels et al., “The nf-core framework for community-curated bioinformatics pipelines,” *Nature Biotechnology*, vol. 38, no. 3, pp. 276–278, 2020.
- [7] H. Li, “Aligning sequence reads, clone sequences and assembly contigs with BWA-MEM,” *arXiv preprint arXiv:1303.3997*, 2013.
- [8] B. Langmead and S. L. Salzberg, “Fast gapped-read alignment with Bowtie 2,” *Nature Methods*, vol. 9, no. 4, pp. 357–359, 2012.